

Integer arithmetic

Arquitectura de Computadores
Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

Luís Nogueira (lmn@isep.ipp.pt)

Today: Integer arithmetic

- Encoding integers
- Addition
- Negation
- Multiplication
- Shifting

Integer arithmetic

- **Many beginning programmers are surprised to find that...**
 - adding two positive numbers can yield a negative result
 - the comparison $x < y$ can yield a different result than the comparison $x - y < 0$
- **These properties are artifacts of the **finite nature of computer arithmetic****
 - As we will see, the integer arithmetic performed by computers is really a form of modular arithmetic
- **Understanding the nuances of computer arithmetic can help programmers write more reliable code**

Today: Integer arithmetic

- Addition
- Negation
- Multiplication
- Shifting

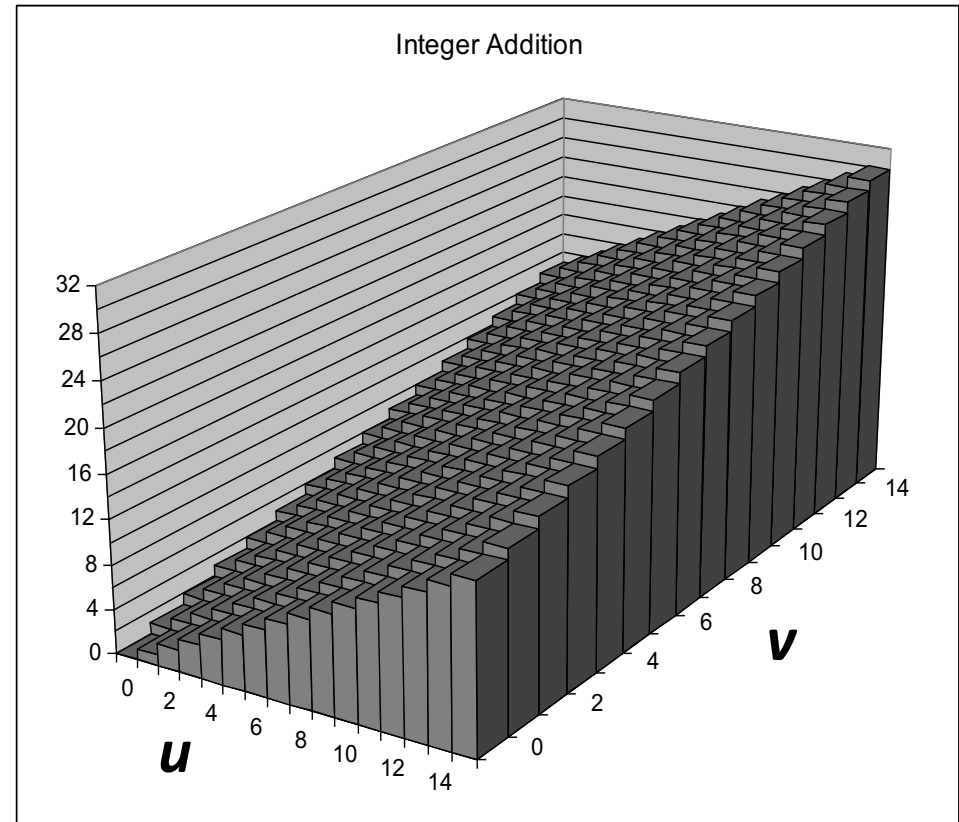
Visualizing (mathematical) integer addition

■ Example:

- 4-bit integers u, v
- Compute true sum with $\text{Add}_4(u, v)$

■ Values increase linearly with u and v

- Forms planar surface



$$\text{Add}_4(u, v) = u + v$$

Unsigned addition in C

Operands: w bits

u 

$+ v$ 

True sum: $w+1$ bits

$u + v$ 

Discard carry: w bits

$\text{UAdd}_w(u, v)$ 

- Unsigned addition ignores carry output

- Implements **modular arithmetic**

$$s = \text{UAdd}_w(u, v) = (u + v) \bmod 2^w$$



$$\text{UAdd}_w(u, v) = \begin{cases} u + v, & u + v < 2^w \\ u + v - 2^w, & u + v \geq 2^w \end{cases}$$

Visualizing unsigned addition

■ Wraps around

- If true sum $\geq 2^w$
- At most once

True Sum



$$\text{UAdd}_4(u, v) = (u + v) \bmod 16$$

Two's complement addition in C

Operands: w bits

u 

$+$ v 

True sum: $w+1$ bits

$u + v$ 

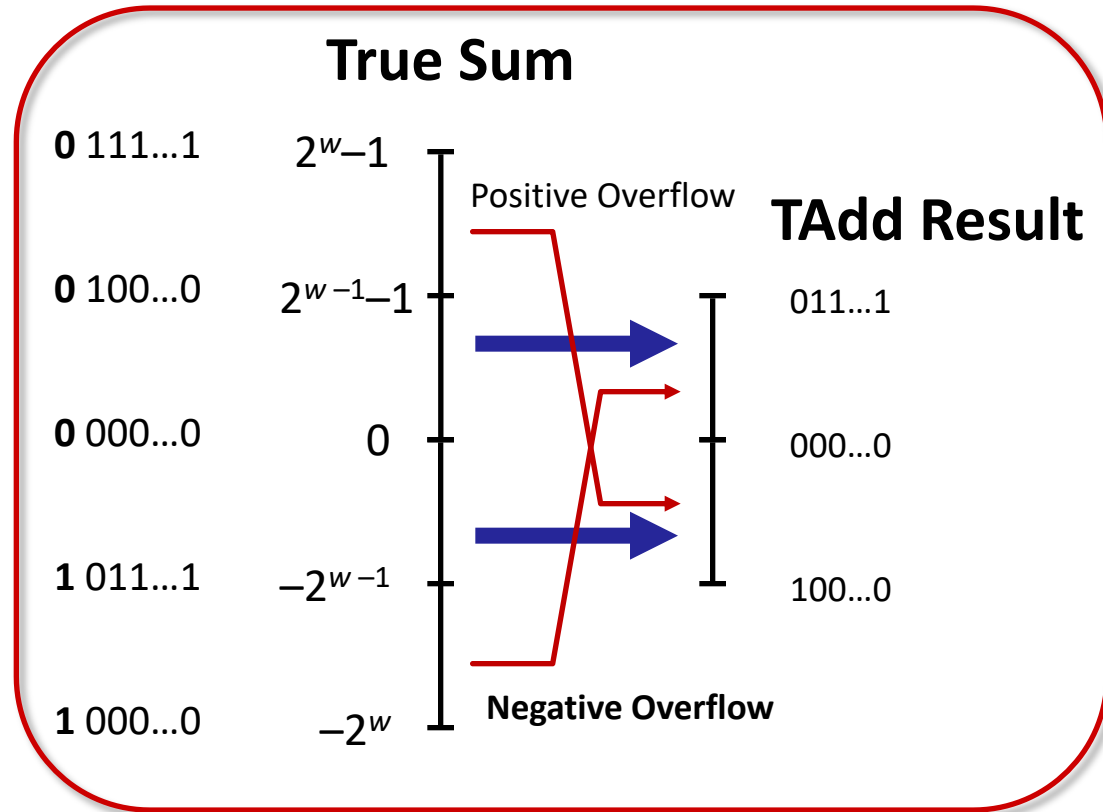
Discard carry: w bits

$\text{TAdd}_w(u, v)$ 

- As before, we avoid ever-expanding data sizes by truncating the representation to w bits
- However, we must decide what to do when the result is either too large (positive) or too small (negative) to represent

Two's complement addition overflow

- Drop off most significant bit
- Treat remaining bits as **two's complement integer**



$$TAdd_w = \begin{cases} u + v + 2^{w-1}, & u + v < TMin_w & \text{(Negative Overflow)} \\ u + v, & TMin_w \leq u + v \leq TMax_w \\ u + v - 2^{w-1}, & TMax_w < u + v & \text{(Positive Overflow)} \end{cases}$$

Visualizing two's complement addition


■ Example:

- 4-bit two's complement
- Range from -8 to +7

■ Wraps Around

- If $\text{sum} \geq 2^{w-1}$
 - Becomes negative
 - At most once
- If $\text{sum} < -2^{w-1}$
 - Becomes positive
 - At most once

Negative Overflow



$$\text{TAdd}_4(u, v) = \text{U2T}_4[(x + y) \bmod 16]$$

Practice problem

- Write a function with the following prototype:

```
int add_ok(int x, int y);
```

- This function should return 1 if arguments x and y can be added without causing overflow

Practice problem

```
int add_ok(int x, int y) {  
    int sum = x+y;  
    int neg_over = x < 0 && y < 0 && sum >= 0;  
    int pos_over = x >= 0 && y >= 0 && sum < 0;  
  
    return !neg_over && !pos_over;  
}
```

- In C, overflows are not signaled as errors
- We can check if an overflow has occurred on $x + y$ by seeing, if and only if,
 - $\text{sum} < x$ (or equivalently, $\text{sum} < y$) for unsigned addition
 - The above conditions for signed addition

Summary: Signed vs Unsigned addition

- **TAdd and UAdd have identical bit-level behavior**
 - Most CPUs use the same machine instruction to perform either unsigned or signed addition
- **Signed vs. Unsigned addition in C**

```
int s, t, u, v;
```

```
s = (int) ((unsigned) u + (unsigned) v);
```

```
t = u + v
```

Will give s == t

Today: Integer arithmetic

- Addition
- **Negation**
- Multiplication
- Shifting

Two's complement negation

- Find the corresponding negative/positive number with the same absolute value
- Every number x in the range $-2^{w-1} \leq x < 2^{w-1}$ has an additive inverse
 - For $x \neq -2^{w-1}$, we can see that its additive inverse is simply $-x$
- For $x = -2^{w-1} = \text{TMin}_w$, $-x = 2^w$ which cannot be represented as a w -bit number

$$x = \begin{cases} -2^{w-1}, & x = -2^{w-1} \\ -x, & x > -2^{w-1} \end{cases}$$

Two's complement negation

- Solution: complement and increment

$$\sim x + 1 == -x$$

- Observation:

$$\sim x + x == 1111\dots111 == -1$$

x	1	0	0	1	1	1	0	1
+ $\sim x$	0	1	1	0	0	0	1	0
-1	1	1	1	1	1	1	1	1

Examples: Complement & Increment

x = 15213

	Decimal	Hex	Binary
x	15 213	3B 6D	00111011 01101101
~x	-15 214	C4 92	11000100 10010010
~x+1	-15 213	C4 93	11000100 10010011

x = 0

	Decimal	Hex	Binary
0	0	00 00	00000000 00000000
~0	-1	FF FF	11111111 11111111
~0+1	0	00 00	00000000 00000000

Today: Integer arithmetic

- Addition
- Negation
- **Multiplication**
- Shifting

Multiplication

- **Computing product of w -bit numbers x and y**
 - Either signed or unsigned
- **Exact results can require as many as $2*w$ bits to represent**
 - Most cases would fit into $2*w - 1$ bits, but the special case of 2^{2w-2} requires the full $2*w$ bits (to include a sign bit of 0)
- **So, maintaining exact results...**
 - Would need to keep expanding word size with each product computed
 - Is done in software, if needed (e.g., by “arbitrary precision” arithmetic packages)

Unsigned multiplication in C

Operands: w bits

u

$*$ v

True product: $2 \cdot w$ bits

$u \cdot v$

Truncate: w bits

$\text{UMult}_w(u, v)$

■ Standard multiplication algorithm

- Truncate result to w -bit number

■ Implements modular arithmetic

$$\text{UMult}_w(u, v) = (u \cdot v) \bmod 2^w$$

Two's complement multiplication in C

Operands: w bits

u

$*$ v

True product: $2*w$ bits

$u \cdot v$

Truncate: w bits

$\text{TMult}_w(u, v)$

- **Compute exact product and ignores high order bits**
 - Some of which are different for signed vs. unsigned multiplication
- **Lower bits are the same of unsigned multiplication**
 - Treat them as a **two's complement integer**

Signed vs Unsigned multiplication

- **TMul and UMul have identical bit-level behavior for the low-order w bits, even though the full $2w$ -bit differ**
 - Separate instructions are provided in IA32 for signed and unsigned multiplication
- **Multiplication in C is performed by truncating the $2w$ -bit product to w bits**

```
int x, y;  
unsigned ux = (unsigned) x;  
unsigned uy = (unsigned) y;  
  
unsigned up = ux * uy;  
int p = x * y;
```

up == (unsigned) p

Today: Integer arithmetic

- Addition
- Negation
- Multiplication
- **Shifting**

Multiplying by constants

■ Most machines shift and add faster than multiply

- The integer multiply instruction is fairly slow, requiring 10 or more clock cycles
- Other integer operations—such as addition, subtraction, bit-level operations, and shifting—require only 1 clock cycle

■ Compiler optimization

- Replace multiplications by constant factors with combinations of shift and addition/subtraction operations

■ Examples


- $u \ll 3 == u * 8$
- $(u \ll 5) - (u \ll 3) == u * 24$

Power-of-2 multiply with shift

Operands: w bits

True product: $w+k$ bits

Truncate: w bits



■ Operation

- $u \ll k$ gives $u * 2^k$

■ Yield the same result

- Both signed and unsigned
- Even in overflow

Compiled multiplication code

C function

```
long int mul_12(long int x){  
    return x*12;  
}
```

Compiled arithmetic operations

```
shll $3, %eax  
shll $2, %ebx  
addl %ebx, %eax
```

Explanation


```
t1 <- x * 8;  
t2 <- x * 4;  
return t1 + t2;
```

- **C compiler automatically generates shift/add code when multiplying by constant**
 - Always uses logical shift

Power-of-2 divide with shift

- **Integer division on most machines is even slower than integer multiplication**
 - Requiring 30 or more clock cycles
- **Dividing by a power of 2 can also be performed using shift operations**
 - Uses a right shift rather than a left shift
- **Unsigned vs Two's complement**
 - Unsigned: logical shift
 - Two's complement: arithmetic shift

Unsigned power-of-2 divide with shift



■ Quotient of unsigned by power-of-2

- $u \gg k$ gives $\lfloor u / 2^k \rfloor$

■ Uses logical shift

Unsigned power-of-2 divide with shift

	Division	Computed	Hex	Binary
x	15 213	15 213	3B 6D	00111011 01101101
x >> 1	7 606.5	7 606	1D B6	00011101 10110110
x >> 4	950.8125	950	03 B6	00000011 10110110
x >> 8	59.4257813	59	00 3B	00000000 00111011

- **The result of shifting consistently rounds toward zero**
 - As is the convention for integer division

Compiled unsigned division code

C function

```
unsigned long int udiv_8(unsigned long int x){  
    return x/8;  
}
```


Compiled arithmetic operations

```
shrl $3, %eax
```

Explanation

```
# Logical shift  
return x >> 3;
```

Signed power-of-2 divide with shift



■ Quotient of signed by power-of-2

- $x \gg k$ gives $\lfloor x / 2^k \rfloor$

■ Uses **arithmetic** shift

Signed power-of-2 divide with shift

	Division	Computed	Hex	Binary
y	-15 213	-15 213	C4 93	11000100 10010011
y >> 1	-7 606.5	-7 607	E2 49	11100010 01001001
y >> 4	-950.8125	-951	FC 49	11111100 01001001
y >> 8	-59.4257813	-60	FF C4	11111111 11000100


- **For a negative number, arithmetic right shift rounds down rather than toward zero**
 - For a positive number, we have 0 as the most significant bit, and so the effect is the same as for a logical right shift

Correct power-of-2 divide with shift

- We can correct for this improper rounding by “biasing” the value before shifting
- Quotient of negative number by power-of-2
 - Want $\lceil x / 2^k \rceil$ (Round toward 0 – convention for integer division)
 - Compute as $\lfloor (x + 2^k - 1) / 2^k \rfloor$
 - In C: $(x + (1 \ll k) - 1) \gg k$
- This technique exploits the property that $\lceil x/y \rceil = \lfloor (x + y - 1) / y \rfloor$ for integers x and y such that $y > 0$

Correct power-of-2 divide with shift

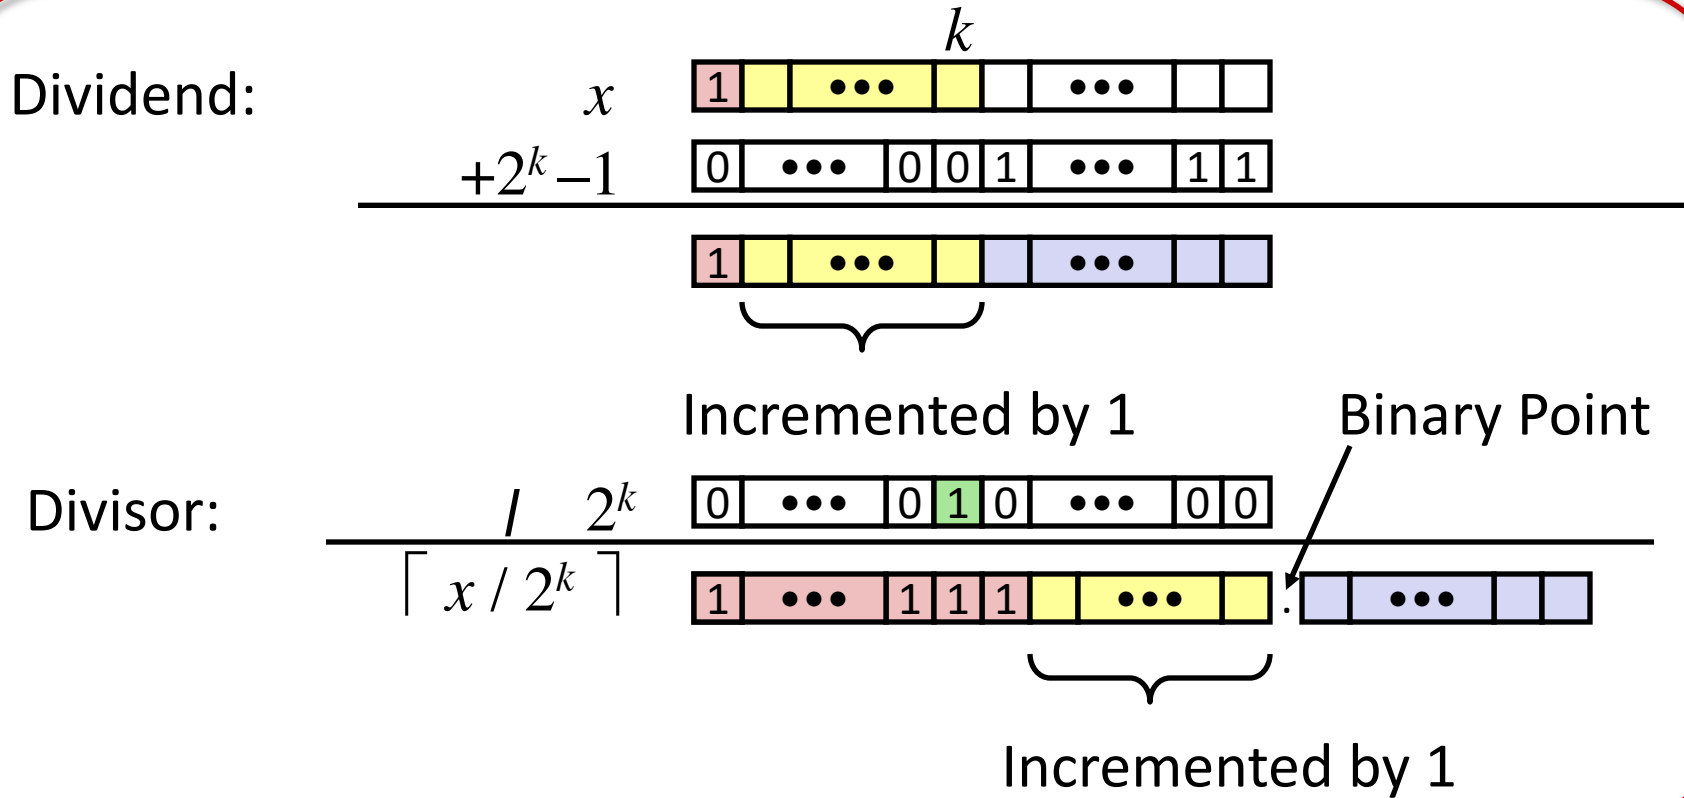
Case 1: No rounding is required ($x \geq 0$)



Biasing has no effect

Correct power-of-2 divide with shift

Case 2: Rounding is required ($x < 0$)



Biasing adds 1 to final result

Compiled signed division code

C function

```
long int idiv8(long int x){  
    return x/8;  
}
```

Compiled arithmetic operations

```
    testl %eax, %eax  
    js    L4  
L3:  
    sarl  $3, %eax  
    ret  
L4:  
    addl  $7, %eax  
    jmp   L3
```

Explanation

```
if x < 0  
    x += 7;  
# Arithmetic shift  
return x >> 3;
```

Arithmetic: Basic rules

■ Addition

- **Unsigned/signed**: Normal addition followed by truncate, same operation on bit level
- **Unsigned**: addition mod 2^w
 - Mathematical addition + possible subtraction of 2^w
- **Signed**: modified addition mod 2^w (result in proper range)
 - Mathematical addition + possible addition or subtraction of 2^w
- **CPUs can use the same machine instruction to perform either unsigned or signed addition**

Arithmetic: Basic rules

■ Multiplication

- **Unsigned/signed**: Normal multiplication followed by truncate, same operation on bit level
- **Unsigned**: multiplication mod 2^w
- **Signed**: modified multiplication mod 2^w (result in proper range)

■ On most machines, integer multiplication and division are fairly slow

- Replace them by constant factors with combinations of shift and addition/subtraction operations

Arithmetic: Basic rules

■ Left shift

- **Unsigned/signed**: multiplication by 2^k
- **Always logical shift**

■ Right shift

- **Unsigned**: logical shift, div (division + round to zero) by 2^k
- **Signed**: arithmetic shift
 - Positive numbers: div (division + round to zero) by 2^k
 - Negative numbers: div (division + round away from zero) by 2^k
 - Use biasing to fix