

Bits and Bytes

Arquitectura de Computadores
Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto


Luís Nogueira (lmn@isep.ipp.pt)

Today: Bits and Bytes

- Representing information as bits
- Bit-level manipulations
- Representations in memory

Everything is bits

- **Binary digits (*bits*) form the basis of the digital revolution**
 - Each bit is 0 or 1
- **Why bits? Electronic implementation**
 - Easy to store with bit elements
 - Reliably transmitted on noisy and inaccurate wires



Everything is bits

- In isolation, a single bit is not very useful
- But, by **encoding/interpreting** sets of bits in various ways
 - Computers determine what to do (instructions)
 - ... and represent and manipulate numbers, sets, strings, etc...
- Most computers use blocks of eight bits, or **bytes**, as the smallest addressable unit of memory

Encoding byte values

- Binary: 00000000₍₂₎ to 11111111₍₂₎
- Decimal: 0₍₁₀₎ to 255₍₁₀₎
- Hexadecimal: 00₍₁₆₎ to FF₍₁₆₎
 - Base 16 number representation
 - Use characters '0' to '9' and 'A' to 'F'

15213: 0011 1011 0110 1101
 3 B 6 D

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Today: Bits and Bytes

- Representing information as bits
- Bit-level manipulations
- Representations in memory

Using bits to represent numbers

- Just like decimal except there are only two digits: 0 and 1
- Everything is based on powers of 2 (1, 2, 4, 8, 16, 32, ...)
- Instead of powers of 10 (1, 10, 100, 1000, ...)
- Binary numbers are shorthands for sums of powers of 2
 - $11011_{(2)} = 1 \times 2^4 + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$
 $= 1 \times 16 + 1 \times 8 + 0 \times 4 + 1 \times 2 + 1 \times 1$
 $= 27_{(10)}$
 - $11101101101101_{(2)} = 15213_{(10)}$
 - $1.1101101101101_{(2)} \times 2^{13} = 1.5213 \times 10^4_{(10)}$

Data representations

- Computers use a **limited number of bits** to encode a number
- Hence, some **operations can overflow/underflow** when the results are too large to be represented

C data type	Intel IA32	x86-64
char	1	1
short int	2	2
int	4	4
long int	4	8
long long int	8	8
float	4	4
double	8	8
long double	10/12	10/16
pointer	4	8

Other types of data

- **Digital representation means that everything is represented by numbers only**
 - Text, code, sound, pictures, ...
- **For sound, pictures, other “real-world” values**
 - Make accurate measurements
 - Convert them to numeric values
- **The usual sequence:**
 - Data is converted into numbers by some mechanism
 - Numbers can be stored, retrieved, processed, transmitted
 - Numbers might be reconstituted into a version of the original

Using bits to represent characters

- **Each character encoded in ASCII format**
 - American Standard Code for Information Interchange
 - Standard 7-bit encoding of character set
- **Each value between 0 and 127 represents a specific character**
- **Most computers extend the ASCII character set to use the full range of 256 characters available in a byte**
 - The upper 128 characters handle special things like accented characters

Using bits to represent characters

	000	001	010	011	100	101	110	111
0000	NULL	DLE		0	@	P	`	p
0001	SOH	DC1	!	1	A	Q	a	q
0010	STX	DC2	"	2	B	R	b	r
0011	ETX	DC3	#	3	C	S	c	s
0100	EDT	DC4	\$	4	D	T	d	t
0101	ENQ	NAK	%	5	E	U	e	u
0110	ACK	SYN	&	6	F	V	f	v
0111	BEL	ETB	'	7	G	W	g	w
1000	BS	CAN	(8	H	X	h	x
1001	HT	EM)	9	I	Y	i	y
1010	LF	SUB	*	:	J	Z	j	z
1011	VT	ESC	+	;	K	[k	{
1100	FF	FS	,	<	L	\	l	
1101	CR	GS	-	=	M]	m	}
1110	SO	RS	.	>	N	^	n	~
1111	SI	US	/	?	O	_	o	DEL

Using bits to represent code

- **A program, from the perspective of the machine, is simply a sequence of bytes**
 - It has no information about the original source program (except some auxiliary tables maintained to aid in debugging)
- **Consider the following C function:**

```
int sum(int x, int y){  
    return x + y;  
}
```

Using bits to represent code

- When compiled on a set of sample machines, we generate machine code having the following byte representations

Linux 32:	55 89 e5 8b 45 0c 03 45 08 c9 c3
Windows:	55 89 e5 8b 45 0c 03 45 08 5d c3
Sun:	81 c3 e0 08 90 02 00 09
Linux 64:	55 48 89 e5 89 7d fc 89 75 f8 03 45 fc c9 c3

- Different machine types use different and incompatible instructions and encodings
 - Even identical processors running different OSes have differences in their coding conventions and hence are not binary compatible

Today: Bits and Bytes

- Representing information as bits
- **Bit-level manipulations**
- Representations in memory

Boolean algebra

■ Developed by George Boole in 19th Century

- Algebraic representation of logic
- Encode “True” as 1 and “False” as 0

And

- $A \& B = 1$ when both $A=1$ and $B=1$

$\&$	0	1
0	0	0
1	0	1

Or

- $A | B = 1$ when either $A=1$ or $B=1$

$ $	0	1
0	0	1
1	1	1

Boolean algebra

■ Developed by George Boole in 19th Century

- Algebraic representation of logic
- Encode “True” as 1 and “False” as 0

Not

- $\sim A = 1$ when $A=0$

\sim	
0	1
1	0

Exclusive-Or (Xor)

- $A \wedge B = 1$ when either $A=1$ or $B=1$, but not both

\wedge	0	1
0	0	1
1	1	0

Extending Boolean algebra

- Operate on bit vectors: operations applied bitwise

01101001	01101001	01101001	
& 01010101	01010101	^ 01010101	~ 01010101
<u> </u>	<u> </u>	<u> </u>	<u> </u>
01000001	01111101	00111100	10101010

- Bitwise operations have many properties in common with integer arithmetic
 - & → Intersection
 - | → Union
 - ^ → Symmetric difference
 - ~ → Complement

Bit-level operations in C

■ Operations `&`, `|`, `~`, `^` available in C

- Apply to any “integral” data type (signed and unsigned)
 - long, int, short, char, ...
- View arguments as bit vectors
- Arguments applied bit-wise

Examples (char data type)

<code>~0x41₍₁₆₎</code>		<code>→ 0xBE₍₁₆₎</code>
<code>~01000001₍₂₎</code>		<code>→ 10111110₍₂₎</code>
<code>~0x00₍₁₆₎</code>		<code>→ 0xFF₍₁₆₎</code>
<code>~00000000₍₂₎</code>		<code>→ 11111111₍₂₎</code>
<code>0x69₍₁₆₎</code>	<code>& 0x55₍₁₆₎</code>	<code>→ 0x41₍₁₆₎</code>
<code>01101001₍₂₎</code>	<code>& 01010101₍₂₎</code>	<code>→ 01000001₍₂₎</code>
<code>0x69₍₁₆₎</code>	<code> 0x55₍₁₆₎</code>	<code>→ 0x7D₍₁₆₎</code>
<code>01101001₍₂₎</code>	<code> 01010101₍₂₎</code>	<code>→ 01111101₍₂₎</code>

Contrast: Logic operations in C

■ Contrast to logical operators `&&`, `||`, `!`

- View 0 as “False”
- Anything nonzero as “True”
- Always return 0 or 1
- **Early termination**

Watch out for `&&` vs. `&` (and `||` vs. `|`)...
a common mistake in C programming

Examples (char data type)


<code>!0x41</code>	<code>→</code>	<code>0x00</code>	<code>(false)</code>
<code>!0x00</code>	<code>→</code>	<code>0x01</code>	<code>(true)</code>
<code>!!0x41</code>	<code>→</code>	<code>0x01</code>	<code>(true)</code>
<code>0x69 && 0x55</code>	<code>→</code>	<code>0x01</code>	<code>(true)</code>
<code>0x69 0x55</code>	<code>→</code>	<code>0x01</code>	<code>(true)</code>
<code>p && *p</code>			<code>(avoids null pointer access)</code>

Bit-level operations

- **One common use of bit-level operations is to implement *masking* operations**
 - A *mask* is a bit pattern that indicates a selected set of bits within a word
- **Very useful in a number of practical applications**
 - IP addressing and subnetting
 - Hash tables
 - Controlling devices
 - Image processing
 - ...

Example: Printer status register

```
#define EMPTY    01
#define JAM      02
#define LOW_INK  16
#define CLEAN    64
```



```
char status;

if (status == (EMPTY | JAM)) ...;
if (status == EMPTY || status == JAM) ...;
while (!(status & LOW_INK)) ...;

status |= CLEAN; /* turns on CLEAN bit */
status &= ~JAM;  /* turns off JAM bit */
```

Shift operations

- C also provides a set of *shift* operations for shifting bit patterns to the left and to the right
- **Arithmetic operators have precedence over shifts**
 - Getting the precedence wrong in C expressions is a common source of program errors, and often these are difficult to spot by inspection
- **For a n-word size value, the shift amount should be a value between 0 and $n - 1$**
 - **Undefined behavior** when shift amount < 0 or \geq word size

Shift operations

■ Left shift: $x \ll y$

- Shift bit-vector x left y positions
 - Throw away extra bits on left
 - Fill with 0's on right


Argument x	01100010
$\ll 3$	00010000
Log. $\gg 2$	00011000
Arith. $\gg 2$	00011000

■ Right shift: $x \gg y$

- Shift bit-vector x right y positions
 - Throw away extra bits on right
- **Logical** shift
 - Fill with 0's on left
- **Arithmetic** shift
 - Replicate most significant bit on left

Argument x	10100010
$\ll 3$	00010000
Log. $\gg 2$	00101000
Arith. $\gg 2$	11101000

Example: Printer status register



```
#define COUNT (8|16|32|64|128)
```


```
/* extract to c */
unsigned int c = (status & COUNT) >> 3;

/* insert v */
status = ((v << 3) | ~COUNT) | (status & ~COUNT);
```


Today: Bits and Bytes

- Representing information as bits
- Bit-level manipulations
- **Representations in memory**

Byte-oriented memory organization



- **Programs refer to data by address**
 - Conceptually, envision it as a very large array of bytes
 - In reality, it's not, but can think of it that way
 - An address is like an index into that array


- **The OS provides private address spaces to each process**
 - Think of a process as a program being executed
 - So, a program can clobber its own data, but not that of others

Machine words

- **Any given computer has a “word size”**
 - Nominal size of integer-valued data and of addresses
- **Until recently, most machines used 32 bits as word size**
 - Limits addresses to 4GB (2^{32} bytes)
- **Increasingly, machines have 64-bit word size**
 - Potentially, could have 18 PB (petabytes) of addressable memory
 - That's 18.4×10^{15}
- **Machines still support multiple data formats**
 - Fractions or multiples of word size
 - Always integral number of bytes

Word-oriented memory organization

- **Addresses specify byte locations**
 - Address of first byte in word
- **Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)**



Pointers in C

- **A *pointer* is a reference to another variable (memory location) in a program**

```
int b = -15213;  
int *p1 = &b;  
char *p2 = (char*)&b;
```

- **The value of a pointer in C is the address of the first byte of some block of storage to which it points to**
 - Whether it points to an integer, a structure, or some other program object

Pointers in C

- **Pointer type impacts pointer arithmetic and the number of bytes that are written/read to/from memory**
 - The computed value is scaled according to the size of the data type referenced by the pointer

```
int b = -15213;  
int *p1 = &b;           // p1+1 increments 4 bytes  
char *p2 = (char*)&b;    // p2+1 increments 1 byte
```

- **Pointer size is determined by the the full word size of the machine**
 - 4 bytes in IA32, 8 bytes in x86-64


Byte ordering

- So, how are the bytes within a multi-byte word ordered in memory?
- **Big Endian convention**
 - Least significant byte has highest address
 - Adopted by Sun, PPC Mac, **Internet**
- **Little Endian convention**
 - Least significant byte has lowest address
 - Adopted by **x86**, ARM processors running Android, iOS, and Windows


Byte ordering example

- Assume `int x` has value of **0x01234567**...
- ... and the address given by `&x` is **0x100**

Big Endian



Little Endian




Representing integers

Decimal: 15213


Binary: 0011 1011 0110 1101

Hex: 3 B 6 D


`int A = 15213;`



`long int C = 15213;`



`int B = -15213;`




Two's complement representation

Representing arrays

```
int A[6] = {1,2,3,4,5,6};
```


- A fixed-size sequential collection of elements of the same type
- The size of the block depends on the number and type of elements
- **Contiguous memory locations**
 - The lowest address corresponds to the first element
 - The highest address corresponds to the last element
 - Elements can be accessed by indexing the array name



Representing strings

```
char S[6] = "18213";
```

- Represented by array of characters
- Each character encoded in ASCII format
 - Character '0' has code 0x30
 - Digit i has code $0x30 + i$
- String should be null-terminated
 - Final character is 0
- Compatibility
 - Byte ordering is not an issue



Examining data representations

■ Code to print byte representation of data

- Casting pointer to *unsigned char** allows treatment as a byte array

```
void show_bytes(unsigned char* start, int len){  
    int i;  
    for (i = 0; i < len; i++)  
        printf("%p\t0x%.2x\n", start+i, *(start+i));  
    printf("\n");  
}
```

***printf* directives:**

%p: Print pointer

%x: Print hexadecimal

Example: `show_bytes` execution

```
int a = 15213; /* 0x3B6D */  
show_bytes((unsigned char*) &a, sizeof(int));
```

Output (Linux/IA32):

0x7fffb7f71dbc	6d
0x7fffb7f71dbd	3b
0x7fffb7f71dbe	00
0x7fffb7f71dbf	00

Summary

- **Computers encode information as bits, generally organized as sequences of bytes**
- **Different models of computers use different conventions for encoding numbers and for ordering the bytes within multi-byte data**
- **The C language is designed to accommodate a wide range of word sizes and numeric encodings**
 - Understanding these encodings at the bit level is important for writing programs that operate correctly over the full range of numeric values