

Integers

Arquitectura de Computadores
Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

Luís Nogueira (lmn@isep.ipp.pt)

Today: Integers

- Encoding integers
- Conversion, casting
- Expanding, truncating
- When should you use unsigned?
- Integer arithmetic

Introduction

- Integers represent a **growing and underestimated source of vulnerabilities** in C and C++ programs
- Integer range checking has not been systematically applied in the development of most C and C++ software
 - Security flaws involving integers exist
 - A portion of these are likely to be vulnerabilities
- A software vulnerability may result when a program evaluates an integer to an unexpected value
 - It is no longer acceptable to assume a program will operate normally given a range of expected inputs when an attacker is looking for input values that produce an abnormal effect

Encoding integers

- **C supports a variety of *integral* data types**
 - `char, short, int, long, long long`
 - Along with an indication of whether the represented numbers are all nonnegative (declared as `unsigned`), or possibly negative (the default)
- **The most common computer representation of signed numbers is known as **two's-complement** form**
- **In two's complement, most significant bit indicates sign**
 - 0 for non-negative
 - 1 for negative

Encoding integers

- **Unsigned**

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

- **Signed: Two's Complement**

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

Sign bit



Two's complement: Simple example ($w=5$)

 $10 =$

-16	8	4	2	1
0	1	0	1	0

$$8+2 = 10$$

 $-10 =$

-16	8	4	2	1
1	0	1	1	0

$$-16+4+2 = -10$$

Encoding integers: another example

```
/* 00111011 01101101(2) */
short x = 15213;

/* 11000100 10010011(2) */
short y = -15213;
```

Weight	15213	-15213
1	1	1
2	0	1
4	1	0
8	1	0
16	0	1
32	1	0
64	1	0
128	0	1
256	1	0
512	1	0
1024	0	1
2048	1	0
4096	1	0
8192	1	0
16384	0	1
-32768	0	1
Sum	15213	-15213

Numeric ranges

■ Unsigned values

- $UMin = 0$
 $= 000\dots 0$
- $UMax = 2^w - 1$
 $= 111\dots 1$

■ Two's Complement values

- $TMin = -2^{w-1}$
 $= 100\dots 0$
- $TMax = 2^{w-1} - 1$
 $= 011\dots 1$
- $-1 = 111\dots 1$

Example: values for $w = 16$ (short)

	Decimal	Hex	Binary
UMax	65 535	FF FF	11111111 11111111
TMax	32 767	7F FF	01111111 11111111
TMin	-32 768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Values for different word sizes

	W			
	8	16	32	64
UMax	255	65 535	4 294 967 295	18 446 744 073 709 551 615
TMax	127	32 767	2 147 483 647	9 223 372 036 854 775 807
TMin	-128	-32 768	-2 147 483 648	-9 223 372 036 854 775 808

■ Observations

- **Asymmetric range**
 - $|TMin| = TMax + 1$
- $UMax = 2 * TMax + 1$

■ C Programming

- `#include <limits.h>`
- Declares constants, e.g.,
 - `ULONG_MAX`
 - `LONG_MAX`
 - `INT_MIN`
- Values are platform specific

Today: Integers

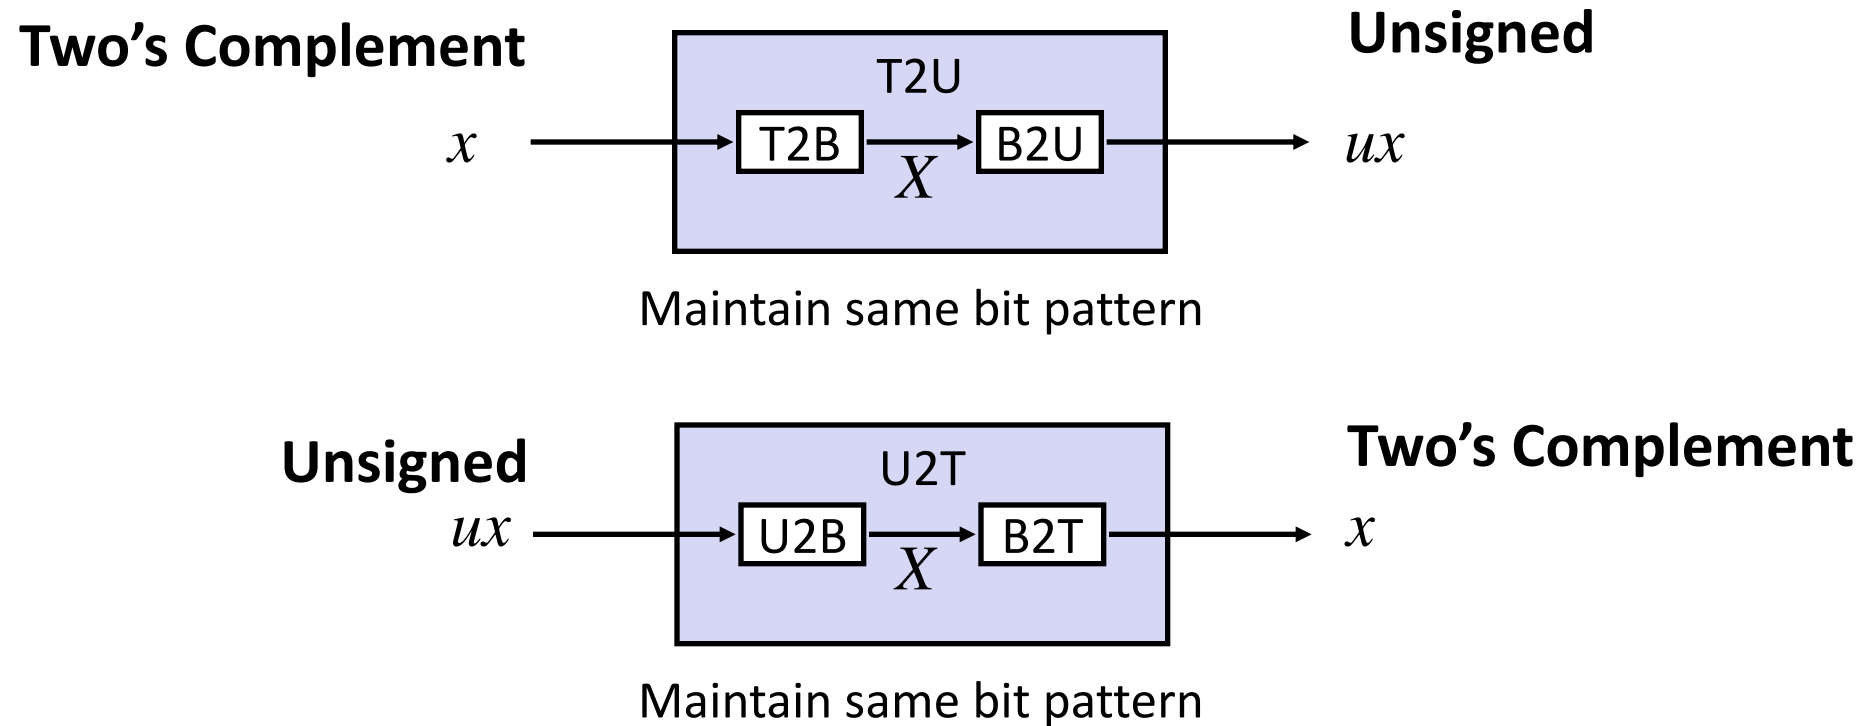
- Encoding integers
- **Conversion, casting**
- Expanding, truncating
- When should you use unsigned?
- Integer arithmetic

Unsigned & signed numeric values

X	$U(X)$	$T(X)$
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

- Same encodings for non-negative values
- Every bit pattern represents an unique integer value
- Each representable integer has an unique bit encoding



Mapping signed \leftrightarrow unsigned



Mappings between unsigned and two's complement numbers:
Keep bit representations and reinterpret

Relation between signed & unsigned

Two's Complement




**Large negative weight
becomes
large positive weight**

Conversion visualized

■ Two's Complement → Unsigned

- Ordering Inversion
- Negative → Big Positive



Casting: Signed vs. unsigned in C

■ Constants

- By default are considered to be signed integers
- Unsigned if have “U” as suffix
 - Examples: `0U`, `4294967259U`

■ Explicit casting between *signed* and *unsigned*

- Same as U2T and T2U

```
int tx, ty;  
unsigned int ux, uy;  
  
tx = (int)ux;  
uy = (unsigned int)ty;
```


Casting: Signed vs. unsigned in C

■ Implicit casting

- Occurs via assignments and procedure calls

```
int tx, ty;  
unsigned int ux, uy;  
  
tx = ux;           /* Cast to signed */  
uy = ty;           /* Cast to unsigned */  
uy = fun(ty);      /* int fun(unsigned int x); */
```

■ Expression evaluation

- If there is a mix of unsigned and signed in a single expression, ***signed values implicitly cast to unsigned***
- Including comparison operations <, >, ==, <=, >=

Casting surprises

TMin = -2147483648

TMax = 2147483647

UMax = 4294967295

UMin = 0

■ Examples for $w = 32$

Constant ₁	Constant ₂	Relation	Evaluation
0	0U	==	unsigned
-1	0	<	signed
-1	0U	>	unsigned
2147483647	-2147483648	>	signed
2147483647U	-2147483648	<	unsigned
-1	-2	>	signed
(unsigned)-1	-2	>	unsigned
2147483647	2147483648U	<	unsigned
2147483647	(int)2147483648U	>	signed

Summary

Casting signed \leftrightarrow unsigned: Basic rules

- Bit pattern is maintained but reinterpreted
- Can have unexpected effects: adding or subtracting 2^w
- Expression containing *signed* and *unsigned* values \rightarrow
signed cast to unsigned

Today: Integers

- Representation: Signed and Unsigned
- Conversion, casting
- **Expanding, truncating**
- When should you use unsigned?
- Integer arithmetic


Expanding bit representation

■ Task:

- Given w -bit signed integer x
- Convert it to $(w+k)$ -bit integer **with same value**


■ Rule:

- Make k copies of sign bit:
- $X' = \underbrace{x_{w-1}, \dots, x_{w-1}}_{k \text{ copies of MSB}}, x_{w-1}, x_{w-2}, \dots, x_0$




Sign extension: Simple example ($w = 5$)

Positive number



Negative number



Expanding bit representation: Example

```

short x = 15213;
int  ix = (int)x;  /* expand to 32-bit value */

short y = -15213;
int  iy = (int)y;  /* expand to 32-bit value */

```

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
ix	15213	00 00 3B 6D	00000000 00000000 00111011 01101101
y	-15213	C4 93	11000100 10010011
iy	-15213	FF FF C4 93	11111111 11111111 11000100 10010011

Truncating numbers

■ Task:


- Given w -bit signed/unsigned integer x
- Convert it to k -bit integer ($k < w$), **possibly changing its value**

■ Rule:

- Drop the high-order $w-k$ bits:

- $X' = \underbrace{x_{k-1}, x_{k-2}, \dots, x_0}_{w-k \text{ low-order bits}}$

- **Reinterpret number**



Truncation: Simple example

No sign change

	-16	8	4	2	1
2 =	0	0	0	1	0

	-8	4	2	1
2 =	0	0	1	0

	-16	8	4	2	1
-6 =	1	1	0	1	0

	-8	4	2	1
-6 =	1	0	1	0

Sign change

	-16	8	4	2	1
10 =	0	1	0	1	0

	-8	4	2	1
-6 =	1	0	1	0

	-16	8	4	2	1
-10 =	1	0	1	1	0

	-8	4	2	1
6 =	0	1	1	0

Truncating numbers: Example

```
int    x    = 53191;

short sx = (short)x;      /* truncate to 16-bit value */

int    y    = sx;        /* expand to 32-bit value */
```

	Decimal	Hex	Binary
x	53191	00 00 CF C7	00000000 00000000 11001111 11000111
sx	-12345	CF C7	11001111 11000111
y	-12345	FF FF CF C7	11111111 11111111 11001111 11000111

Relative order of conversion

- Consider the following code:

```
short sx = -12345;          /* 0xcfc7 */
unsigned int uy = sx;       /* Mystery!!! */

printf("uy = %u : \t", uy);
show_bytes((unsigned char*) &uy, sizeof(unsigned int));
```

- When run on a little-endian machine, we get:

```
uy = 4294954951 : c7 cf ff ff
```

- First changes the size and then from *signed* to *unsigned*

- Otherwise, we would get `uy = 53191 : c7 cf 00 00`

Summary

Expanding, Truncating: Basic rules

- **Expanding (e.g., *short* to *int*)**
 - Unsigned: zeros added
 - Signed: sign extension
 - Both yield expected result

- **Truncating (e.g., *unsigned int* to *unsigned short*)**
 - Unsigned/signed: bits are truncated (mod operation)
 - Result is reinterpreted
 - For small numbers yields expected behavior

Today: Integers

- Representation: Signed and Unsigned
- Conversion, casting
- Expanding, truncating
- **When should you use unsigned?**
- Integer arithmetic

When should you use unsigned?

- ***Do not*** use without understanding its implications
- **Easy to make mistakes**

```
unsigned int i;  
int cnt;  
.  
.  
.  
for (i = cnt-2; i >= 0; i--)  
    a[i] += a[i+1];
```

- **Can be very subtle**

```
#define DELTA sizeof(int)  
int i, cnt;  
.  
.  
.  
for (i = cnt; i-DELTA >= 0; i-= DELTA)  
    .  
    .  
    .
```

Counting down with unsigned

■ Proper way to use *unsigned* as loop index

```
unsigned int i;  
int cnt;  
  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

■ Even better

```
size_t i;  
  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

- Data type **size_t** defined as *unsigned* value with length equal to word size
- Code will work even if **cnt** = *UMax*
- What if **cnt** is signed and < 0?

When should you use unsigned? (cont.)

- **Do use when performing modular arithmetic**
 - Multi-precision arithmetic, in which numbers are represented by arrays of words
- **Do use when using bits to represent sets**
 - Logical right shift, no sign extension
- **Addresses are naturally unsigned, so system programmers find unsigned types to be helpful**
- **See Robert Seacord, “*Secure Coding in C and C++*”**

Today: Integers

- Representation: Signed and Unsigned
- Conversion, casting
- Expanding, truncating
- When should you use unsigned?
- **Integer arithmetic**

Integer arithmetic

- **Many beginning programmers are surprised to find that...**
 - adding two positive numbers can yield a negative result
 - the comparison $x < y$ can yield a different result than the comparison $x - y < 0$
- **These properties are artifacts of the **finite nature of computer arithmetic****
 - Integer arithmetic performed by computers is really a form of **modular arithmetic**
 - See “Integer arithmetic” document for details
- **Understanding the nuances of computer arithmetic can help programmers write more reliable code**

Arithmetic: Basic rules

■ Addition

- True value requires **w+1 bits** in the worst case
- **Unsigned/signed**: Normal addition **followed by truncate and reinterpretation**
- **Unsigned**: addition mod 2^w
 - Mathematical addition + possible subtraction of 2^w
- **Signed**: modified addition mod 2^w (result in proper range)
 - Mathematical addition + possible addition or subtraction of 2^w

■ CPUs can use the same machine instruction to perform either unsigned or signed addition/subtraction

- Same operation on bit level

Unsigned vs Signed addition

unsigned char

1110 1001	E9	223
+ 1101 0101	+ D5	+ 213
<hr/>	<hr/>	<hr/>
1 1011 1110	1BE	446
<hr/>	<hr/>	<hr/>
1011 1110	BE	190

char

1110 1001	E9	-23
+ 1101 0101	+ D5	+ -43
<hr/>	<hr/>	<hr/>
1 1011 1110	1BE	-66
<hr/>	<hr/>	<hr/>
1011 1110	BE	-66

Arithmetic: Basic rules

■ Multiplication

- True value requires **$2*w$ bits** in the worst case
- **Unsigned/signed**: Normal multiplication **followed by truncate and reinterpretation**
- **Unsigned**: multiplication mod 2^w
- **Signed**: modified multiplication mod 2^w (result in proper range)

■ On most machines, integer multiplication and division are fairly slow

- Separate instructions are provided in IA32 for signed and unsigned multiplication/division
- Replace them by constant factors with combinations of shift and addition/subtraction operations

Unsigned vs Signed multiplication

unsigned char

	1110 1001	E9	233
*	1101 0101	* D5	* 213
	<u>1100 0001 1101 1101</u>	<u>C1DD</u>	<u>49629</u>
	1101 1101	DD	221

char

	1110 1001	E9	-23
*	1101 0101	* D5	* -43
	<u>0000 0011 1101 1101</u>	<u>03DD</u>	<u>989</u>
	1101 1101	DD	-35

Arithmetic: Basic rules

■ Left shift

- **Unsigned/signed**: multiplication by 2^k
- **Always logical shift**

■ Right shift

- **Unsigned**: logical shift, div (division + round to zero) by 2^k
- **Signed**: arithmetic shift
 - Positive numbers: div (division + round to zero) by 2^k
 - Negative numbers: div (division + round away from zero) by 2^k
 - Use biasing to fix

C Puzzle

■ For each of the following C expressions, either:

- Argue that is true for all argument values
- Give example where not true

Initialization

```
int x = foo();
int y = bar();
unsigned ux = x;
unsigned uy = y;
```

- $x < 0 \Rightarrow ((x*2) < 0)$
- $ux \geq 0$
- $x \ \& \ 7 == 7 \Rightarrow (x \ll 30) < 0$
- $ux > -1$
- $x > y \Rightarrow -x < -y$
- $x * x \geq 0$
- $x > 0 \ \&\& \ y > 0 \Rightarrow x + y > 0$
- $x \geq 0 \Rightarrow -x \leq 0$
- $x \leq 0 \Rightarrow -x \geq 0$

C Puzzle answers

■ *TMin* makes a good counterexample in many cases

- | | | |
|----------------------|------------------------------|--|
| ■ $x < 0$ | $\Rightarrow ((x*2) < 0)$ | False: <i>TMin</i> |
| ■ $ux \geq 0$ | | True: $0 = UMin$ |
| ■ $x \& 7 == 7$ | $\Rightarrow (x \ll 30) < 0$ | True: $x_1 = 1$ |
| ■ $ux > -1$ | | False: 0 |
| ■ $x > y$ | $\Rightarrow -x < -y$ | False: $-1, TMin$ |
| ■ $x * x \geq 0$ | | False: 65535
$(x*x=-131071)$ |
| ■ $x > 0 \&\& y > 0$ | $\Rightarrow x + y > 0$ | False: <i>TMax, TMax</i> |
| ■ $x \geq 0$ | $\Rightarrow -x \leq 0$ | True: $-TMax < 0$ |
| ■ $x \leq 0$ | $\Rightarrow -x \geq 0$ | False: <i>TMin</i> |

Summary

- **In C there are two different ways bits can be used to encode integers**
 - One that can only represent nonnegative numbers
 - One that can represent negative, zero, and positive numbers
- **Unsigned values can be very useful**
 - But the implicit conversion of signed to unsigned can lead to errors or vulnerabilities
- **It is possible to expand or shrink an encoded integer to fit a representation with a different length**

Summary

- **Integer arithmetic performed by computers is really a form of modular arithmetic**
 - These properties are artifacts of the finite nature of computer arithmetic