

# Computer Architecture (Practical Class)

## *Introduction to the C Programming Language*

*aka C for Java Programmers*

Luís Nogueira    Raquel Faria

Departamento de Engenharia Informática  
Instituto Superior de Engenharia do Porto

`{lmn,arf}@isep.ipp.pt`

2020/2021

# The C programming language

- C is a general-purpose, imperative programming language, developed in the 70es at Bell Labs
- Its history is closely related with the UNIX operating system (OS) as it was created out of the necessity to port the OS to different machines
- It was designed to be simple and small. The book “The C Programming Language, Second Edition” by Brian Kernighan and Dennis Ritchie describes the whole language, the standard library and several examples and exercises with 261 pages
- C is widely used for systems programming (OS kernel, drivers, compilers, ...) and embedded systems
- Examples: Linux OS and MySQL have been written in C

# Why C?

- C ranks amongst the most popular programming languages (1st in the TIOBE Index of 2020, 2nd in 2019)
- C influenced many other programming languages (C++, Java, Objective-C, Swift, C#, PHP, Go, ...)
- Provides access to low-level mechanisms such as memory management, explicit initialization and error detection
- C usually produces very efficient programs

## Philosophical reason

- Helps understanding what happens in the system from the UI to the electrons ;-)

- Operators, conditionals, loops and other languages constructs are similar to Java
- Operators:
  - Arithmetic: `+, -, *, \, %, ++, --, +=, -=, *=, ...`
  - Relational: `<, >, <=, >=, ==, !=`
  - Logical: `&&, ||, !, ? :`
  - Bit: `&, |, ^, ~, <<, >>`
- Language constructs:
  - `if( ){ } else { }`
  - `while( ){ }`
  - `do { } while( )`
  - `for(i=0; i<100; i++){ }`
  - `switch( ) { case 0: ... }`
  - `break, continue, return`
- No exception handling statements

<b>what</b>	<b>C</b>	<b>Java</b>
<b>type of language</b>	function oriented	object oriented
<b>basic programming unit</b>	function	class = ADT
<b>portability of source code</b>	possible with discipline	yes
<b>portability of compiled code</b>	no, recompile for each architecture	yes, bytecode is 'write once, run anywhere'
<b>compilation</b>	creates machine language code	creates Java virtual machine language bytecode
<b>execution</b>	loads and executes program	interprets byte code
<b>variable auto-initialization</b>	not guaranteed	all variables must be initialized; compile-time error to access uninitialized variables

Source: <http://introcs.cs.princeton.edu/java/faq/c2java.html>

# "Hello World" program in C

Listing 1: hello.c (get it: <http://codepad.org/q5x1Quqa>)

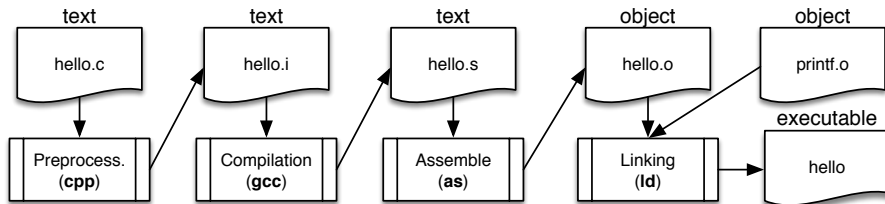
```
/*
our program uses printf(), defined in the C standard library "stdio"
lines starting with '#' are called preprocessor directives, and do not
have a ';' at the end.
(this is a multiline comment)
*/
#include <stdio.h>

/*
the function main() returns an integer and receives no arguments (void)
main() is the first function to be called when the program is executed
*/
int main(void)
{
    /* printf() prints formatted output;
       \n is a newline */
    printf("Hello, World ! \n");

    // the main function returns the value 0 (single line comment)
    return 0;
}
```

# Compiling C programs

- C programs must be transformed into machine-code so they can be executed, in a process called **compilation**
- The compilation process involves several other steps:
  - preprocessing; compilation; assembling; linking



We will use a compiler called the *GNU C Compiler - gcc*

- Compile hello world program: `gcc -Wall hello.c -o hello`
- The gcc option “-save-temps” tells it to keep all files generated during compilation
- Example: `gcc -save-temps hello.c -o hello`

# Important gcc switches

option	Description
-Wall	all warnings – <b>use always!</b>
-o <i>filename</i>	output filename for object or executable
-c	compile only, do not link; used to create an object file (.o) for a single (non-main) .c file (module)
-g	insert debugging information
-E	stop after the preprocessing stage; output goes to standard output
-v	show information about gcc and/or compilation process
-S	performs preprocessing and compilation only; that is, convert C source into assembly
-save-temps	keep temporary files created (.i, .s, .o, ...)
-l <i>library-name</i>	link with library called <i>library-name</i>
-I <i>dir</i>	add <i>dir</i> to the list of dirs to be searched for header files
-L <i>dir</i>	add <i>dir</i> to the list of dirs to be searched for the libraries specified with -l;

More at: <http://goo.gl/exA6gY>



- Always read the output of gcc carefully!
  - Gives pretty good indications about the origin of the error
  - Several sources of errors:
    - preprocessor: missing include files;
    - parser: syntax errors;
    - assembler: syntax errors in assembly code (only if you are coding assembly);
    - linker: missing libraries.
  - Often, one error causes lots of subsequent errors
    - fix first error, and then retry – ignore the rest.
  - Often, errors are caused by previous mistakes
    - for example a missing ';' often will cause an error in the subsequent line(s).

Compile with *-Wall* and *do not* ignore warnings!

- Warnings often provide indication about errors that will manifest themselves at runtime

- The C preprocessor (cpp) allows defining macros, which are brief abbreviations for longer constructs
- Preprocessor directives start with a '#' at the beginning of the line and are used for:
  - Inserting content of another file into file to be compiled: `#include`
  - Conditional compilation: `#if`; `#ifdef`
  - Definition of macros and constants: `#define`.
- Before compilation, the preprocessor reads the source code and transforms it.

- Example 1:

```
#include <stdio.h> // searches for stdio.h in system defined directories
#include "mydefs.h" // searches for mtdefs.h in the current directory
```

- Example 2:

```
#define MAX 100
#define check(x) ((x) < MAX)
if check(i) { ... }
```

- Becomes:

```
if ((i) < 100) { ... }
```

## Use the C preprocessor with caution

- It is easy to introduce subtle errors
- Not visible in debugging
- Code hard to read

## Integer types

Type	Storage size	Value range
char	1 byte	-128 to 127
unsigned char	1 byte	0 to 255
int	4 bytes (IA-32)	-2,147,483,648 to 2,147,483,647
unsigned int	4 bytes (IA-32)	0 to 4,294,967,295
short	2 bytes	-32,768 to 32,767
unsigned short	2 bytes	0 to 65,535
long	4 bytes	-2,147,483,648 to 2,147,483,647
unsigned long	4 bytes	0 to 4,294,967,295

## Floating-point types

Type	Storage size	Value range	Precision
float	4 byte	1.2E-38 to 3.4E+38	6 decimal places
double	8 byte	2.3E-308 to 1.7E+308	15 decimal places
long double	10 byte	3.4E-4932 to 1.1E+4932	19 decimal places

```
char c='A';
```

```
char c=100;
```

```
int i=-2343234;
```

```
unsigned int ui=100000000;
```

```
float pi=3.14;
```

```
double long_pi=0.31415e+1;
```

- The storage size of some types **varies** among architectures. E.g. A *long* is 4 bytes in IA32 and 8 bytes in x86-64 machines
- *char* is misleading. It is a numeric type that happens to be sometimes used to store ASCII character codes
- The *void* type comprises an empty set of values; it is an incomplete type that cannot be completed.
  - You cannot define variables of type *void*, however *void* can be used to:
    - indicate that a function has no parameters. E.g. `int func(void);`
    - indicate that a function has no return. E.g. `void func(int n);`
    - define a pointer that does not specify the type it points to (more on this in the following lectures). E.g. `void* ptr;`
- Two kinds of type conversions
  - Implicit: automatic type conversion by the compiler.  
E.g.: `int a=1000; char b=a; // b=-24 (lower 8 bits of a=11101000...)`
  - Explicit: explicitly defined by the programmer.  
E.g.: `float f=1.2; int d=(int)f; // d=1`

- C has a unary operator `sizeof`, that can be used to get the storage size of variables and data types, *measured in the number of char type storage size*.
- Examples:
  - `sizeof(int)`: returns the size of *int*
  - `sizeof int`: (same as above) returns the size of *int*
  - `sizeof a`: returns the size of the variable *a*
  - `sizeof(char)`: returns the size of type `char`; **guaranteed to always be 1**

### Important

- While, for most modern systems, the `char` type has 8 bits, *there is no guarantee that this is always true*.
  - The number of bits of type `char` is defined in the `CHAR_BIT` constant in `<limits.h>`.
- 
- Check the file `<limits.h>` for the sizes and limits of the integer types
    - e.g. `CHAR_MAX`, `CHAR_MIN`, `INT_MAX`, `INT_MIN`
  - Check the file `<float.h>` for the sizes and limits of the floating-point types
    - e.g. `FLT_MIN`, `FLT_MAX`

- The next slide will present an example using `sizeof` and several constants from `<limits.h>` and `<float.h>`

### `printf()` format specifiers quick reference

- `%d` or `%i`: Signed decimal integer
- `%u`: Unsigned decimal integer
- `%f`: Decimal floating point, lowercase
- `%E`: Scientific notation (mantissa/exponent), uppercase
- `%c`: Character
- `%s`: String of characters
- see: <http://www.cplusplus.com/reference/cstdio/printf/>



Listing 2: sizeof.c (<http://codepad.org/sPFcuyKy>)

```
#include <stdio.h> // needed for printf
#include <limits.h> // needed for CHAR_BIT, INT_MAX, INT_MIN
#include <float.h> // needed for FLT_MAX, FLT_MIN, FLT_DIG

int main() {
    char n='A';

    printf("\nStorage size for variable n: %u\n", sizeof(n));

    printf("\nStorage size for char: %u\n", sizeof(char));
    printf("Number of bits in a char: %u\n", CHAR_BIT);

    printf("\nStorage size for int: %u\n", sizeof(int));
    printf("Minimum int value: %u\n", INT_MIN );
    printf("Maximum int value: %u\n", INT_MAX );

    printf("\nStorage size for float : %d \n", sizeof(float));
    printf("Minimum float positive value: %E\n", FLT_MIN );
    printf("Maximum float positive value: %E\n", FLT_MAX );
    printf("Precision value for float: %d\n", FLT_DIG );

    printf("\nStorage size for double=%u\n", sizeof(double));

    return 0;
}
```

## Using sizeof: Example (3/3)

Output of the example (Listing 2; sizeof.c)

```
Storage size for variable n: 1
```

```
Storage size for char: 1
```

```
Number of bits in a char: 8
```

```
Storage size for int: 4
```

```
Minimum int value: -2147483648
```

```
Maximum int value: 2147483647
```

```
Storage size for float : 4
```


```
Minimum float positive value: 1.175494E-38
```

```
Maximum float positive value: 3.402823E+38
```

```
Precision value for float: 6
```

```
Storage size for double=8
```

- C allows to define arrays of elements of the same type
- Examples of arrays, with statically defined sizes (size is fixed):
  - `int a[10];` // array of 10 integers
  - `int a1[]={1, 2, 3, 4, 5};` // array of 5 integers, initialized to 1, 2, 3, 4 and 5;
  - `int a2[1000]={0};` // array of 1000 integers, all initialized to 0;
  - `short s[100];` // array of 100 shorts
  - `float m[10][10];` // 10x10 matrix of floats
- For an array containing  $N$  elements, indexes are  $0..N - 1$ , accessed using `a[0]`, `a[1]`, ..., `a[N-1]`
- Arrays are stored as a continuous linear arrangement of elements



### gcc **does not check** when you access invalid indexes

- `int x[10]; x[10] = 5;` is an overflow of the array, and will result in undefined behaviour (it may work for a while...usually results in a segmentation fault and program termination)

### An array **cannot be the target of an assignment**

- Assume an array `int v[10]`, declared previously. The following statement is **not valid**: `v = {1, 2, 3, 4, 5};`
- We can only assign values “in bulk” to an array when it is declared. After the declaration, a valid statement would be: `for (i=0; i<5; i++) v[i] = i+1;`
- If you want to copy arrays, use `memcpy(dest, src, size);`

**C does not remember** how large arrays are (i.e., no length attribute)

- `sizeof` works **only** for statically defined arrays, within the scope they are declared

- Example:

```
{  
    int a[10];  
    printf("%u", sizeof(a)); // prints 40 in IA32  
}
```

- The `{}` define the scope of these statements
- Size of array can be computed with `sizeof(a) / sizeof(a[0])`

### Passing the address of an array to a function

- When the array is passed as an argument to a function, the size information is not available

```
void func(int a[ ]) {  
    printf("%u", sizeof(a)); // prints 4 in IA-32  
}
```

- More on this in the following classes...


- The solution is for the programmer to maintain the length of the array:
  - by passing the size as an argument of the function;
  - by using a globally defined constant;
  - by defining a data structure for storing the array and its size together;
  - by defining a value that indicates the end of the the array (e.g. an int array ends with a -1 value).

## C does not have a specific data type for strings

- Strings are just char arrays with a NUL ('\0') terminator (value zero)

```
char a[10]="abc";
```

- Defines an array of 10 chars. The first 3 chars will have the characters 'a', 'b', and 'c'. The fourth will be the NUL terminator:



Listing 3: funcao.c (<http://codepad.org/gZ65gm8F>)

```
int xpto(char s[])
{
    int c=0;

    while (s[c]!=0) c=c+1;

    return c;
}
```

- What is the functionality of the function?
  - A. The function returns the ASCII code of the last character.
  - B. The function returns the number of elements of the string.
  - C. The function returns the number of words of the string.
  - D. None of the above.



## Remember!

- Arrays (strings - arrays of char - included) cannot be target of assignments after the declaration
- The following statement is valid, and declares an array `s[6]`, initialized with the characters 'H', 'e', 'l', 'l', 'o', '\0': `char s[]="Hello";`
- However, after declaring `s[]`, you **cannot assign it a new value**: `s="World";` // this is not a valid statement!
- Keep in mind that strings should be copied with `strncpy(dest_str, src_str, n_chars)`: `strncpy(s, "World", 5);`
- Also, note that it is the programmer's responsibility to check that the destination string has enough storage.

## Example: Compute the average of two integers

Listing 4: avg.c (<http://codepad.org/Jy9uXsjQ>)

```
#include <stdio.h> /* for declaration of printf */

/* globals (really necessary?) */
int n1=6, n2=4, avg=0;

/* this function computes the (integer) average of two integers */
int calc_avg(int a, int b) {
    int c=0; /* local variable */
    c=(a+b)/2;
    return c;
}

/* the program starts by executing this function */
int main(void) {

    avg = calc_avg(n1, n2); /* call function and save return */
    printf("Avg = %d\n", avg);

    return 0; /* returns 0 */
}
```

- Good code should be mostly self-documenting
  - Variables and function names should generally help making clear what you are doing
  - Comments should not describe what the code does, but why. What the code does should be self-evident (assume the reader knows C)
  - Do comment: each source file, function headers, large blocks of code, tricky bits of code (e.g. bit manipulations)
- Use C-style naming conventions:
  - E.g. prefer `get_radius()` to `GetRadius()`;
  - *i* and *j* for loop variables.
- Bodies of functions, loops, if-else statements, etc. should be indented

- Define constants and use them. Constants make your code more readable, and easier to change
- Avoid global variables. Pass variables as arguments to functions
- Initialize variables before using them!
- Use good error detection and handling. *Always* check return values from functions, and handle errors appropriately

## Making the best use of C

- Read <https://goo.gl/5u9aCo> for advice on how to use the C language

- Implement a C program that reads 10 integers into an array and computes their average.
- The average should be calculated in a separate function, but its value should be printed by the main().